What is an anomaly?

ANOMALY DETECTION IN R

Alastair Rushworth
Data Scientist

Defining the term anomaly

Anomaly: a data point or collection of data points that do not follow the same pattern or have the same structure as the rest of the data

Point anomaly

- A single data point

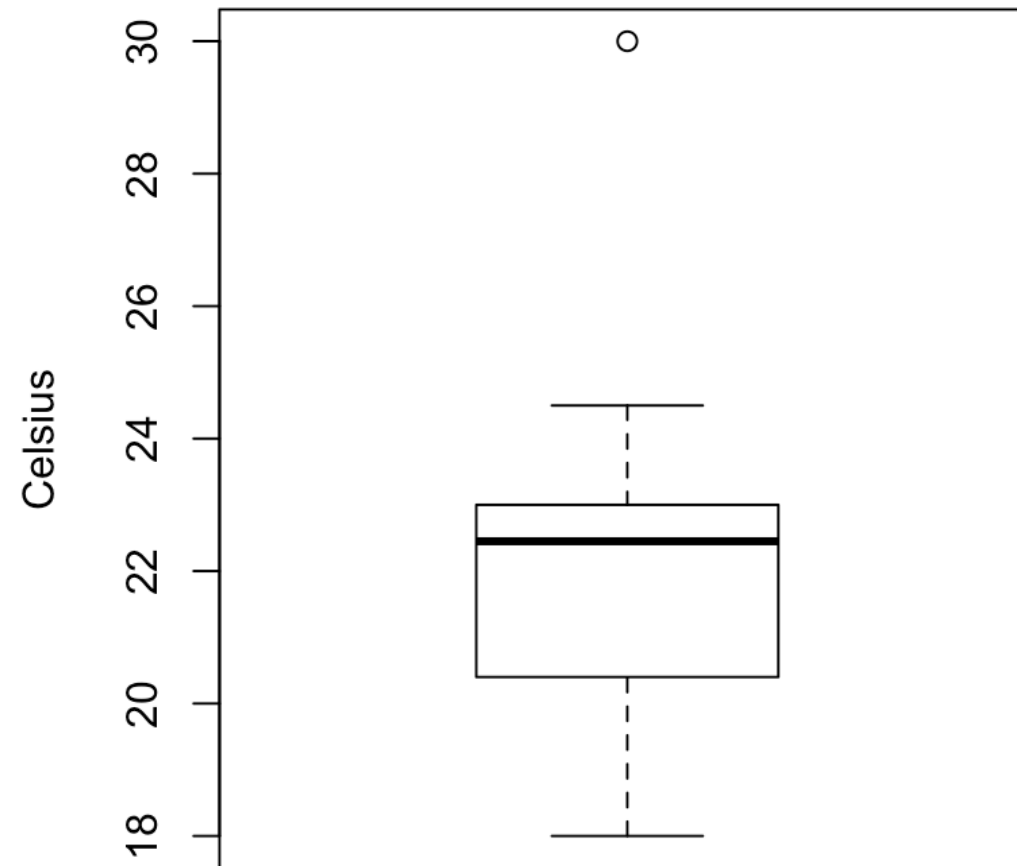- Unusual when compared to the rest of the data

Example: A single 30C daily high temperature among a set of ordinary spring days

```
summary(temperature)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  18.00   20.45   22.45   22.30   22.98   30.00
```
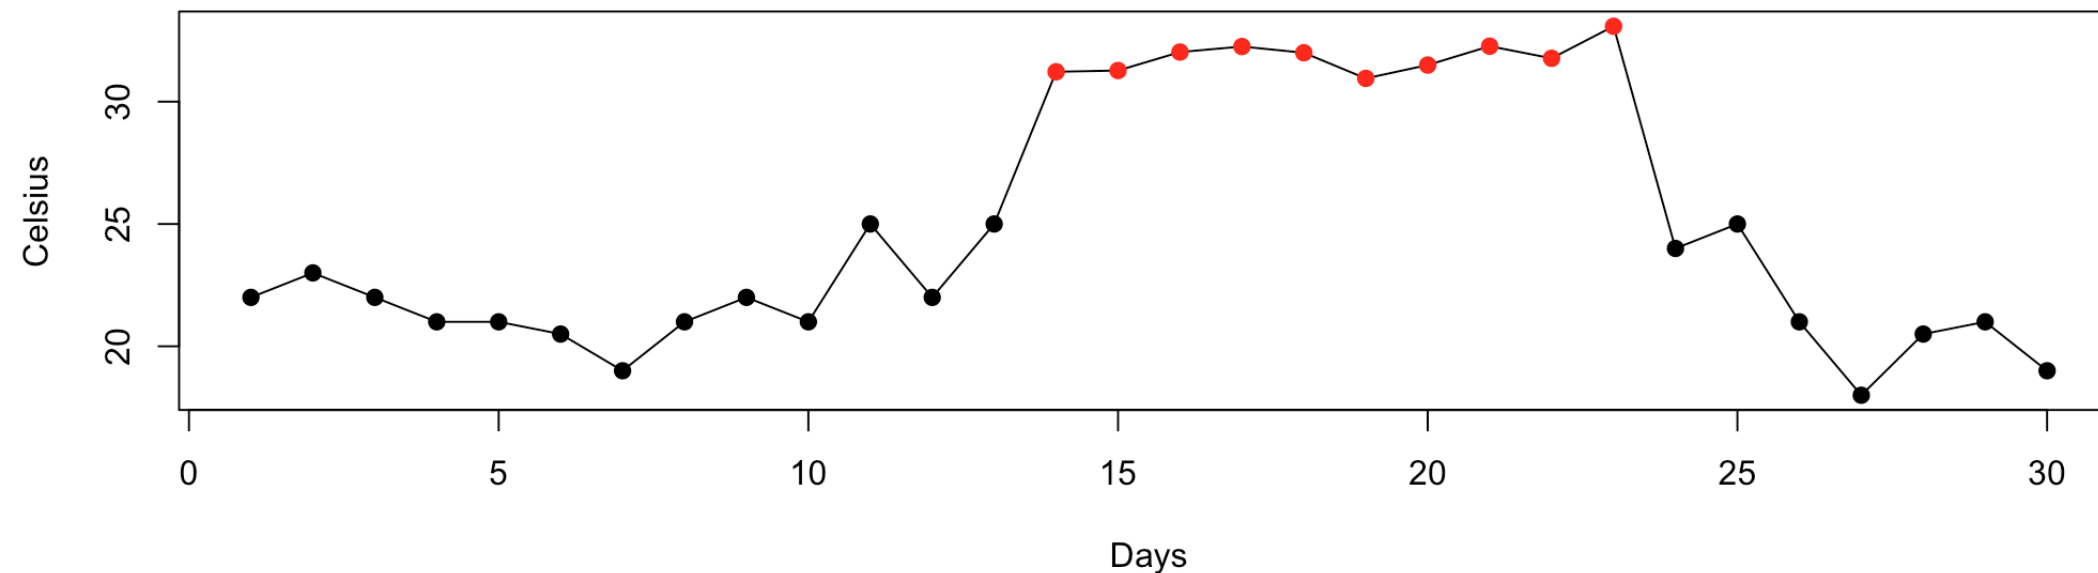
Visualizing point anomalies with a boxplot

```
boxplot(temperature, ylab = "Celsius")
```

Collective anomaly

- An anomalous collection of data instances

- Unusual when considered together

Example: 10 consecutive high daily temperatures

Let's practice!

ANOMALY DETECTION IN R
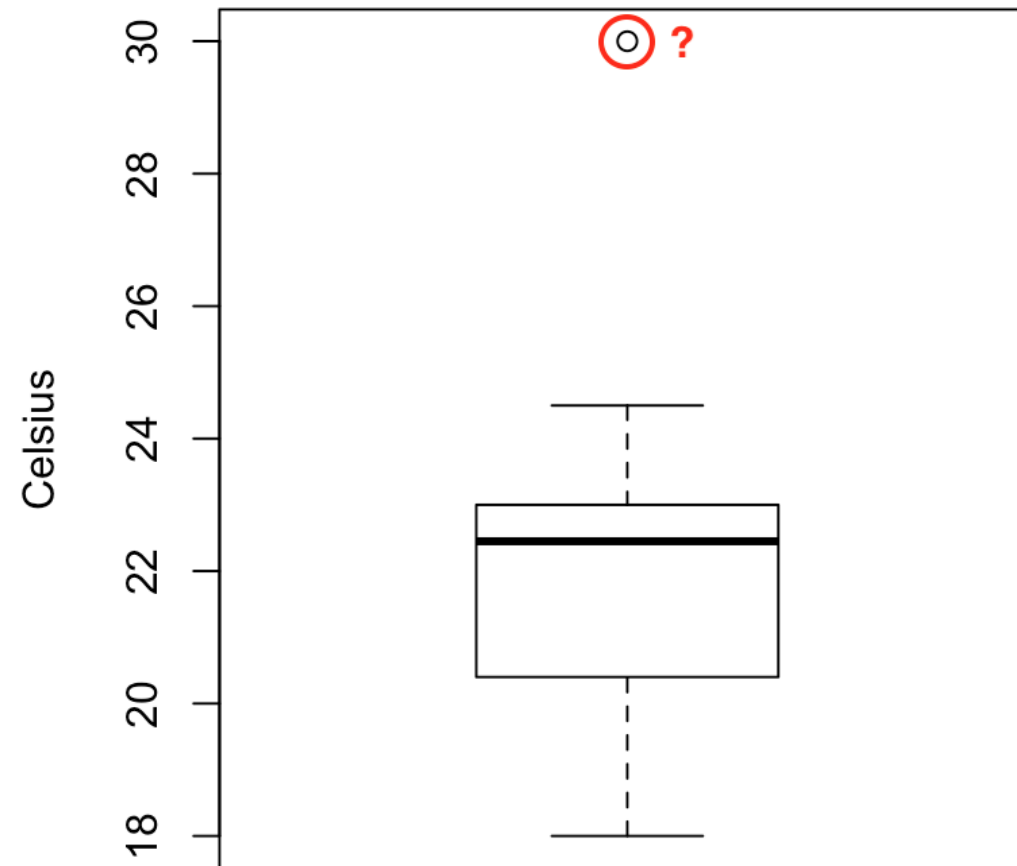
Testing the extremes with Grubbs' test

ANOMALY DETECTION IN R

Alastair Rushworth
Data Scientist

Visual assessment is not always reliable!

```
boxplot(temperature, ylab = "Celsius")
```
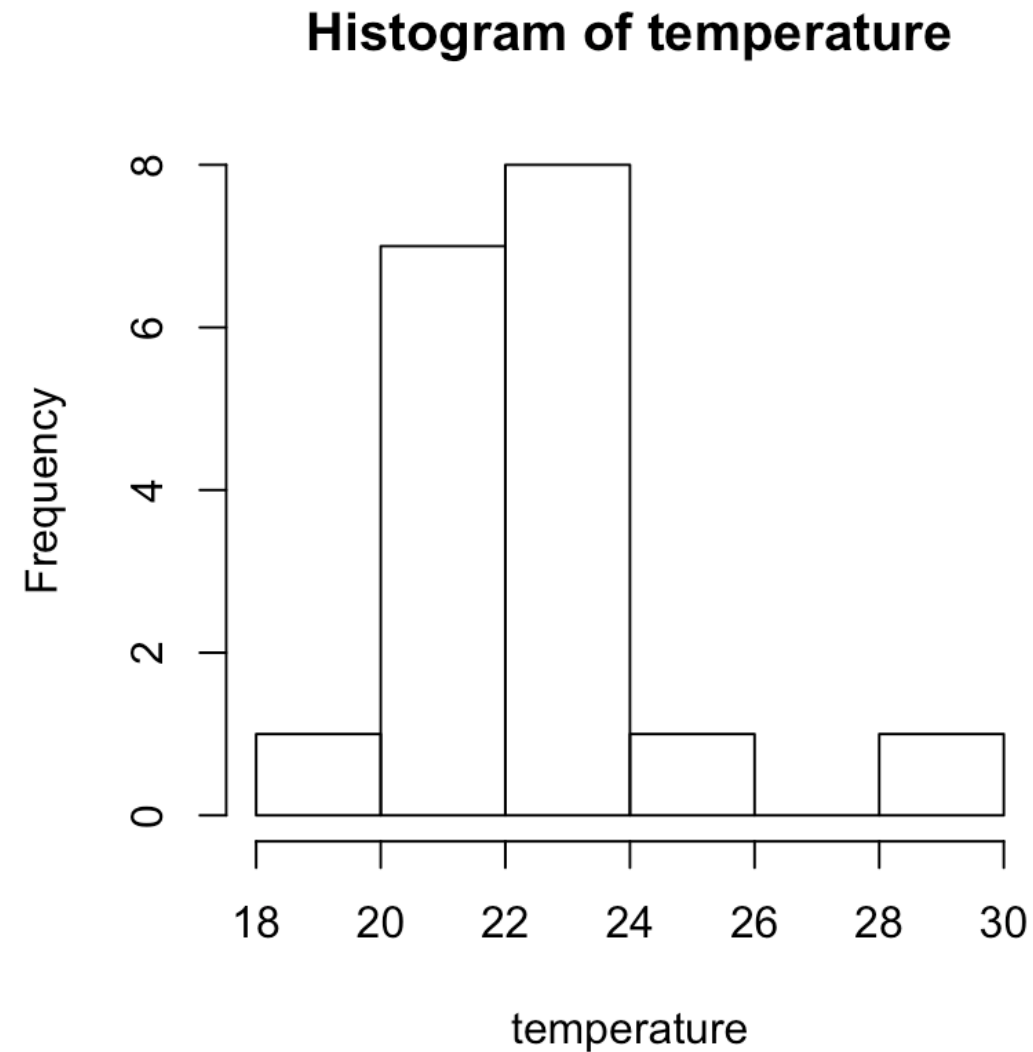
Grubbs' test

- Statistical test to decide if a point is outlying

- Assumes the data are normally distributed

- Requires checking the normality assumption first

Checking normality with a histogram

```
hist(temperature, breaks = 6)
```


Histogram of temperature

Running Grubbs' test

Use the `grubbs.test()` function:

```
grubbs.test(temperature)
```

```
	Grubbs test for one outlier
data:  temp
G = 3.07610, U = 0.41065, p-value = 0.001796
alternative hypothesis: highest value 30 is an outlier
```

Interpreting the p-value

```
grubbs.test(temperature)
```

```
    Grubbs test for one outlier


data:  temperature
G = 3.07610, U = 0.41065, p-value = 0.001796
alternative hypothesis: highest value 30 is an outlier
```

`p-value`

- Near 0 - *stronger* evidence of an outlier

- Near 1 - *weaker* evidence of an outlier

Get the row index of an outlier

Location of the **maximum**

```
which.max(weights)
```

```
5
```

Location of the **minimum**

```
which.min(temperature)
```

```
12
```

Let's practice!

ANOMALY DETECTION IN R

Detecting multiple anomalies in seasonal time series

ANOMALY DETECTION IN R

Alastair Rushworth
Data Scientist

Monthly revenue data
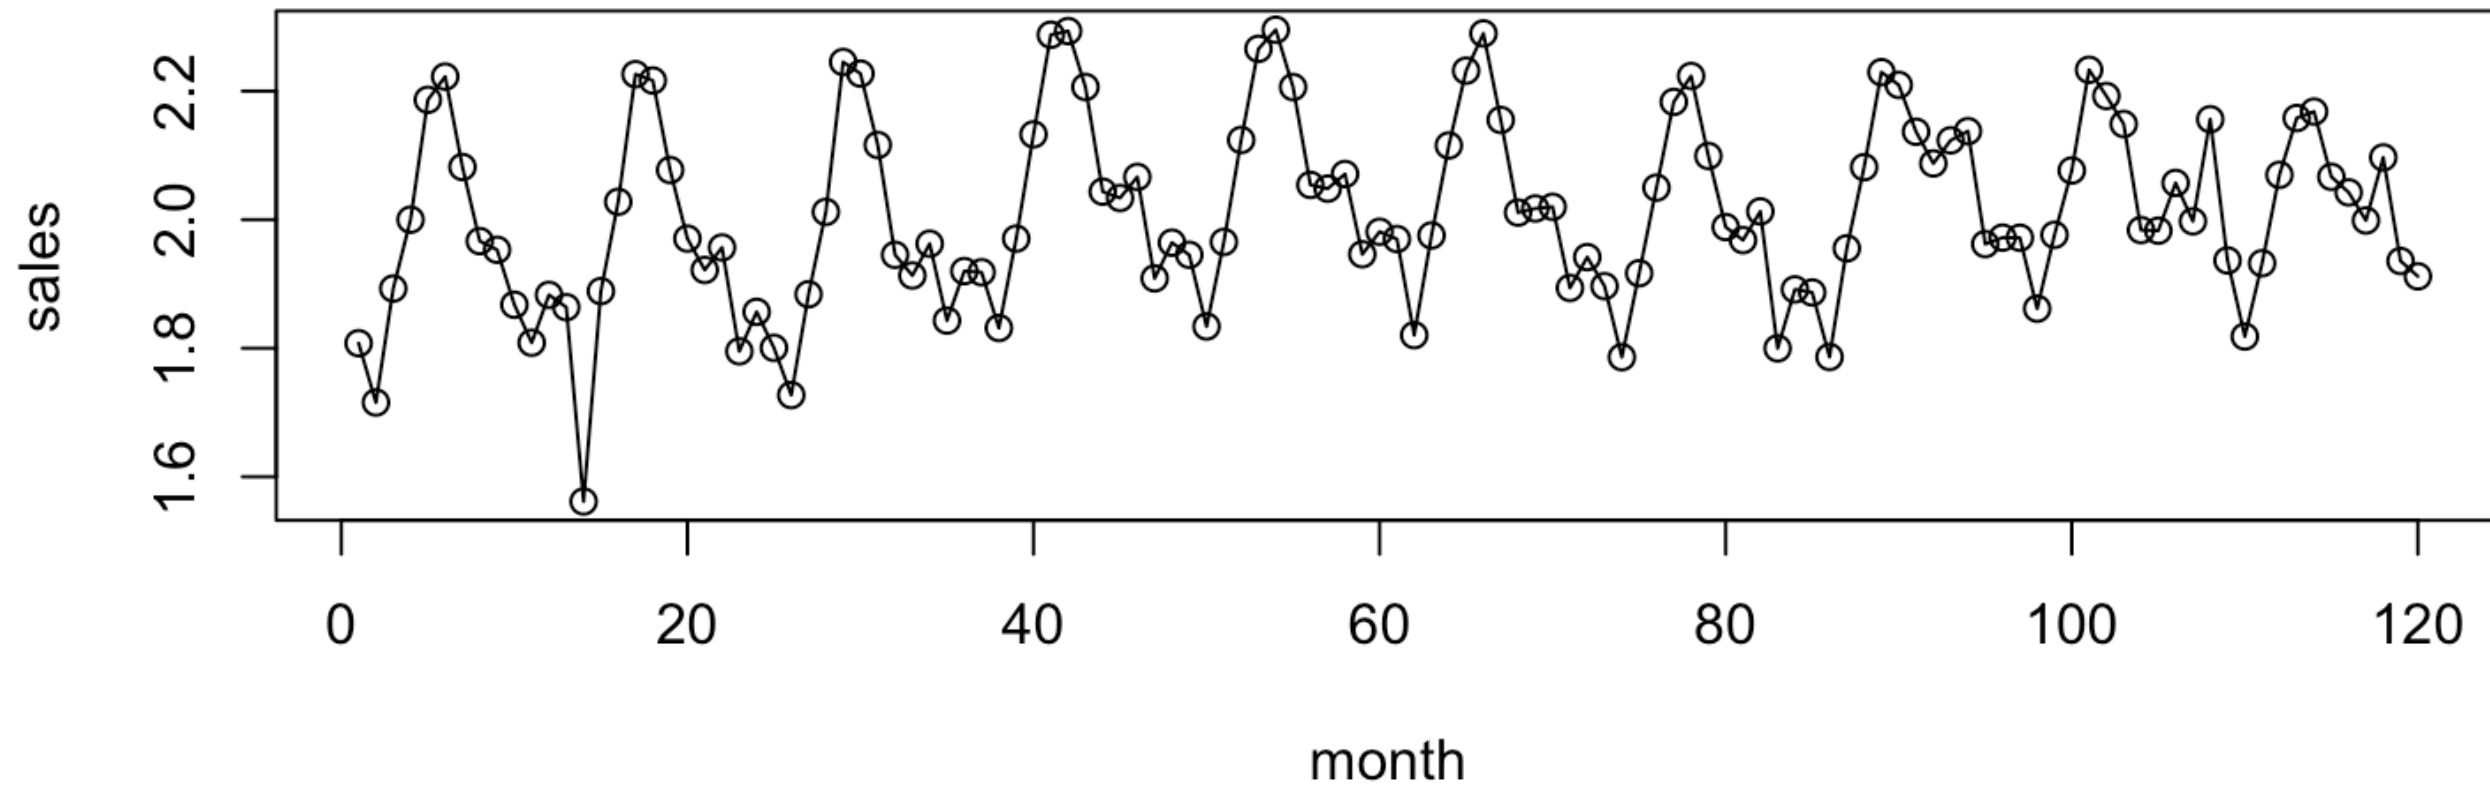
```
head(msales)
```

```
  sales month
1 6.068     1
2 5.966     2
3 6.133     3
4 6.230     4
5 6.407     5
6 6.433     6
```

Grubbs' test not appropriate here

- Seasonality may be present

- May be multiple anomalies

Visualizing monthly revenue

```
plot(sales ~ month, data = msales, type = 'o')
```

Seasonal-Hybrid ESD algorithm usage

```r
library(AnomalyDetection)
sales_ad <- AnomalyDetectionVec(x = msales$sales, period = 12,
                                direction = 'both')
```

Arguments

- `x` : vector of values

- `period` : period of repeating pattern

- `direction` : find anomalies that are small (`'neg'`), large (`'pos'`), or both (`'both'`)

Package download from **https://github.com/twitter/AnomalyDetection**

Seasonal-Hybrid ESD algorithm output
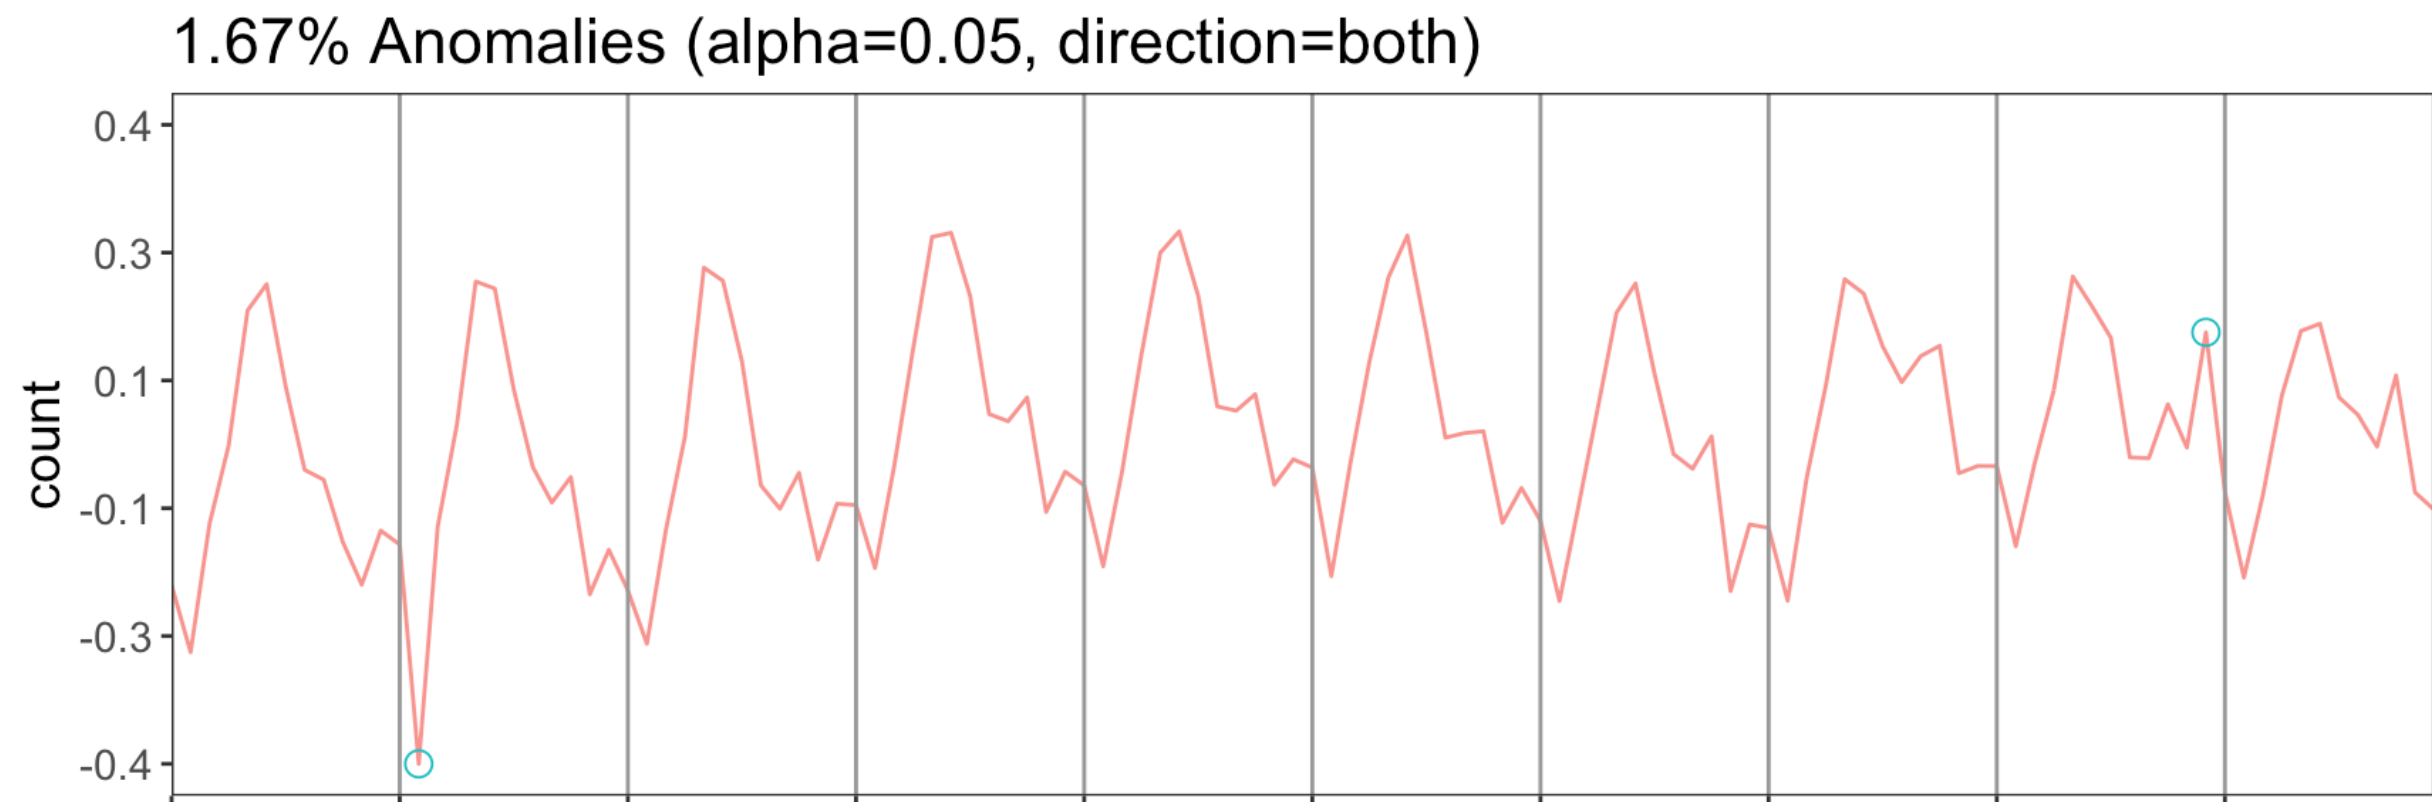
```r
sales_ad <- AnomalyDetectionVec(x = msales$sales, period = 12,
                                direction = 'both')
```

```r
sales_ad$anoms
```

```
  index anoms
1    14 1.561
2   108 2.156
```

Seasonal-Hybrid ESD algorithm plot

```
AnomalyDetectionVec(x = msales$sales, period = 12,
                    direction = 'both', plot = T)
```


1.67% Anomalies (alpha=0.05, direction=both)

Let's practice!